

Convolutional- & Recurrent Neural Networks

Notes for Topics in Social Data Science at UCPH

March 6, 2019

Thor Donsby Noe

1 Convolutional Nets

Issues with vanilla feed-forward neural network

- Data squeezed into a $1 \times N$ vector.
- Single operation on whole input.
- No attention to **spatial adjacency**.

Solution: 'Slide' weight matrix over a part of the input. E.g.

- Input image: $32 \times 32 \times 3$ (i.e. RGB-colors)
- Filter: $5 \times 5 \times 3 \rightarrow \text{dot product} \rightarrow 1$ number
- Activation map: $28 \times 28 \times 1$ (convolved over all spatial locations)
 - One activation map for each filter. These are stacked.
 - E.g. $28 \times 28 \times 6$ for 6 $5 \times 5 \times 3$ filters.
 - * An additional activation map can be added by using a filter on the activation map of the former layer.
 - * E.g. a new activation map layer: $24 \times 24 \times 10$ by using 10 $5 \times 5 \times 3$ filters.

Stride

- How many pixels are moved each time, e.g. 1.

Padding

- The no. zeroes added at the borders of the input image before applying the filter.
- Full padding: If you don't want the image to shrink.
- Half padding: The image shrinks less, e.g. padding = 1.

Width of the activation map will be

$$W = \frac{N + 2P - F}{S} + 1$$

E.g. for the example above with padding=2

$$W = \frac{32 + 2 \cdot 2 - 5}{1} + 1 = 32$$

Pooling: Condense the information by downsampling.

- Speeds up the learning process.
- 'max pool' where only largest value is saved for each quadrant as they are expected to contain the more important signal.
- 'average pooling'
- E.g. $16 \times 16 \rightarrow 4 \times 4$
- Doesn't shrink it in the depth direction.

2 Recurrent Nets

Optimal for **sequential data** but useful for Elements
non-sequential data as well.

Issue for **all** feed forward neural networks

- Input and output must be of hard-assigned dimensions.

Backpropagation: Compute the sum of losses.

- Can get complicated as the number becomes very high up with long chains.
- Can break it into parts.

- f: Forget vector (0 = forget, 1 = remember) → what to forget from c-vector
 - c: New hidden vector, i.e. long term memory
 - c is not converted, → not multiplied by weight → gradient doesn't explode → much quicker
- i, g: adds to memory vector, c
- o: output vector